



# American International University-Bangladesh (AIUB)

Faculty of Science & Technology (FST)

Department of Computer Science

Introduction to Data Science

Mid-Term Project Report

Summer 2024-2025

Section:A

Group:4

SL #	Student Name	Student ID	Contribution (%)
1.	Md Saniul Islam Sony	22-46400-1	100%
2.	Tanvir Ahmed Tuhin	22-46475-1	100%
3.	MD.Ashraful Islam	22-46479-1	100%

## Dataset Description

A credit-risk dataset where each row is one borrower and one loan application. It combines basic demographics, finances, and loan details, plus two binary outcome fields.

All columns:

**person\_age** — borrower's age (years)

**person\_gender** — gender

**person\_education** — highest education level

**person\_income** — annual income

**person\_emp\_exp** — years of employment experience

**person\_home\_ownership** — housing status (e.g., RENT, OWN, MORTGAGE)

**loan\_amnt** — requested loan amount

**loan\_intent** — purpose of the loan (e.g., PERSONAL, MEDICAL)

**loan\_int\_rate** — interest rate (%)

**loan\_percent\_income** — loan payment as a share of income (0–1)

**cb\_person\_cred\_hist\_length** — credit history length (years)

**credit\_score** — credit score (approx. 300–850)

**previous\_loan\_defaults\_on\_file** — whether the borrower ever defaulted before (0/1)

**loan\_status** — outcome of this loan application or performance (0/1)

## 1. Handle missing value:

### Description:

- **Drop missing rows (complete cases).**

We first list rows that contain any NA to inspect what's missing, then create a complete-case view by removing those rows.

- **Replace with mean (numeric).**

For a numeric column that is roughly symmetric (here, `person_income`), we fill NA values with the column mean calculated from the available (non-missing) values.

- **Replace with median (numeric).**

For a numeric column that may be skewed or affected by outliers (here, `person_age`), we fill NA values with the median, which is more robust than the mean.

- **Replace with mode (categorical).**

For a categorical/label column (here, `loan_status`), we fill NA values with the mode (the most frequent category). The `mlv(..., method="mfv")` function returns that most frequent value, which we use to impute the missing entries.

### Code:

- Drop missing value:

```
missing_rows <- data[!complete.cases(data), ]
head(missing_rows)
```

### Output:

	person_age	person_gender	person_education	person_income	person_emp_exp	person_home_ownership	loan_amt	loan_intent	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	credit_score	previous_loan_defaults_on_file	loan_status
1	21	Female	Master	71948	0	RENT	35000	PERSONAL	16.02	0.49	3	561	No	1
2	21	Female	High School	12282	0	OWN	1000	EDUCATION	11.14	0.49	2	504	Yes	0
3	25	Female	High School	12438	3	MORTGAGE	5500	MEDICAL	12.87	0.49	3	635	No	1
4	23	Female	Bachelor	79753	0	RENT	35000	MEDICAL	15.23	0.44	2	672	No	1
5	24	male	Master	66135	1	RENT	35000	MEDICAL	14.27	0.53	4	586	No	1
10	21	female	High School	12739	0	OWN	1600	VENTURE	14.74	0.13	3	640	No	1

- Replace with mean

```
df_mean <- data
df_mean$person_income[is.na(df_mean$person_income)] <-
mean(df_mean$person_income, na.rm=TRUE)
head(df_mean)
```

Output:

```
person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948      0      RENT      35000      PERSONAL      16.02      0.49      3      561      No      1
2      21      female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      Yes      0
3      25      female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      No      1
4      23      female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      No      1
5      24      male      Master      66135      1      RENTY      35000      MEDICAL      14.27      0.53      4      586      No      1
6      23      female      High School      12951      0      OWN      2500      VENTURE      7.14      0.19      2      532      No      1
>
```

- Replace with median

```
df_median <- df_mean#replace by median
df_median$person_age[is.na(df_median$person_age)] <-
median(df_median$person_age, na.rm = TRUE);
head(df_median)
```

Output:

```
person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948      0      RENT      35000      PERSONAL      16.02      0.49      3      561      No      1
2      21      female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      Yes      0
3      25      female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      No      1
4      23      female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      No      1
5      24      male      Master      66135      1      RENTY      35000      MEDICAL      14.27      0.53      4      586      No      1
6      23      female      High School      12951      0      OWN      2500      VENTURE      7.14      0.19      2      532      No      1
>
```

- Replace with mode

```
mode_val <- mlv(df_mode$loan_status, method = "mfv", na.rm = TRUE);
df_mode$loan_status[is.na(df_mode$loan_status)] <- mode_val
head(as.data.frame(df_mode))
```

Output:

```
person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948      0      RENT      35000      PERSONAL      16.02      0.49      3      561      No      1
2      21      female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      Yes      0
3      25      female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      No      1
4      23      female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      No      1
5      24      male      Master      66135      1      RENTY      35000      MEDICAL      14.27      0.53      4      586      No      1
6      23      female      High School      12951      0      OWN      2500      VENTURE      7.14      0.19      2      532      No      1
>
```

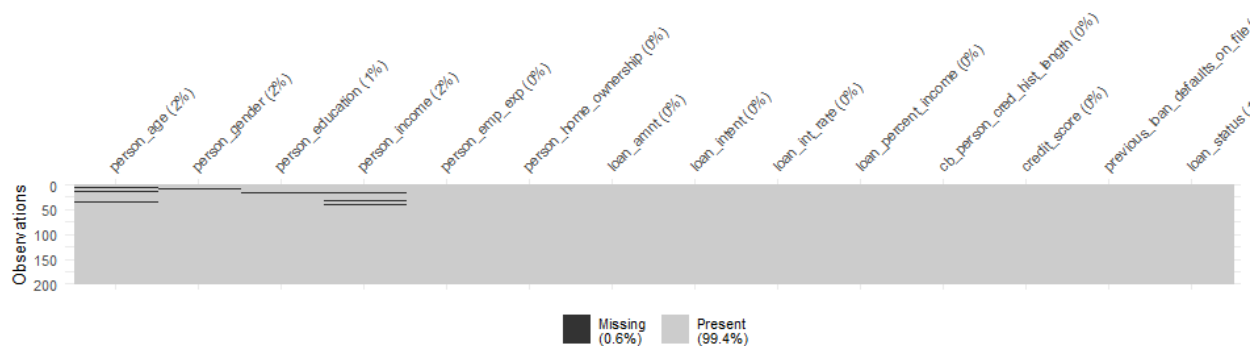
## 2. see missing values on a graph:

### Description:

In this task we will use the `vis_miss()` function from the *naniar* package to visualize missing values in the dataset with a graph. The graph will display variables as columns and records as rows, where shaded blocks indicate missing entries. This

makes it easy to see which columns contain missing data and how much is missing, giving us a clear picture of the dataset's quality.

```
vis_miss(data)
```



### 3. convert attributes from numeric to categorical or categorical to numeric:

#### Description:

In this task we will convert attributes from numeric to categorical or from categorical to numeric depending on how we need them for analysis. For example, the column `loan_status` is stored as numbers (0/1), but for easier interpretation we convert it into labels "Yes" and "No". On the other hand, the column `previous_loan_defaults_on_file` may be stored as text ("Yes"/"No"), so we convert it back into numeric values (1/0) for modeling.

We solve this by using simple `ifelse()` statements:

```
df_mode$loan_status <- ifelse(df_mode$loan_status == 1  
, "Yes", "No");  
head(df_mode)
```

**Output:**

	person_age	person_gender	person_education	person_income	person_emp_exp	person_home_ownership	loan_amt	loan_intent	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	credit_score	previous_loan_defaults_on_file	loan_status
1	21	female	Master	71944	0	RENT	35000	PERSONAL	16.02	0.49	3	561	0	No
2	21	female	High School	12282	0	OWN	1000	EDUCATION	11.14	0.49	2	504	1	No
3	25	female	High School	12438	3	MORTGAGE	5500	MEDICAL	12.87	0.49	3	635	0	No
4	23	female	Bachelor	79753	0	RENT	35000	MEDICAL	15.23	0.44	2	675	0	No
5	24	male	Master	66135	1	RENT	35000	MEDICAL	14.27	0.53	4	586	0	No
6	23	female	High School	12951	0	OWN	2500	VENTURE	7.14	0.19	2	532	0	No

```
df_mode$previous_loan_defaults_on_file <-
  ifelse(df_mode$previous_loan_defaults_on_file == "Yes",
    1,0);
head(df_mode)
```

**Output:**

	person_age	person_gender	person_education	person_income	person_emp_exp	person_home_ownership	loan_amt	loan_intent	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	credit_score	previous_loan_defaults_on_file	loan_status
1	21	Female	Master	71944	0	RENT	35000	PERSONAL	16.02	0.49	3	561	0	No
2	21	Female	High School	12282	0	OWN	1000	EDUCATION	11.14	0.49	2	504	0	No
3	25	Female	High School	12438	3	MORTGAGE	5500	MEDICAL	12.87	0.49	3	635	0	No
4	23	Female	Bachelor	79753	0	RENT	35000	MEDICAL	15.23	0.44	2	675	0	No
5	24	male	Master	66135	1	RENT	35000	MEDICAL	14.27	0.53	4	586	0	No
6	23	Female	High School	12951	0	OWN	2500	VENTURE	7.14	0.19	2	532	0	No

#### 4. Detect outliers in the data set and use the appropriate approach to handle those values:

##### Description:

In this task we detect and handle extreme values in numeric columns using the Interquartile Range (IQR) method. First, we calculate Q1 (25th percentile), Q3 (75th percentile), and the IQR ( $Q3 - Q1$ ). Any value smaller than  $Q1 - 1.5 \times IQR$  is treated as a lower outlier, and any value greater than  $Q3 + 1.5 \times IQR$  is treated as an upper outlier.

The `get_outliers()` function lists these values so we can see which records fall outside the normal range. Then the `fix_outliers()` function corrects them:

- If a value is less than the lower bound, it is replaced by the lower bound.
- If a value is greater than the upper bound, it is replaced by the upper bound.

This way, the dataset does not lose any rows, but all values are adjusted to stay within reasonable limits, reducing the influence of extreme numbers on further analysis.

```
get_outliers <- function(data, colname) {  
  if (!colname %in% names(data)) stop("Column not found.")  
  if (!is.numeric(data[[colname]])) stop("Column must be  
numeric.")  
  
  Q1 <- quantile(data[[colname]], 0.25, na.rm = TRUE)  
  Q3 <- quantile(data[[colname]], 0.75, na.rm = TRUE)  
  IQR <- Q3 - Q1  
  
  lower <- Q1 - 1.5 * IQR  
  upper <- Q3 + 1.5 * IQR  
  
  # return only the outlier values from dataset  
  outliers <- data[[colname]][data[[colname]] < lower |  
data[[colname]] > upper]  
  
  return(outliers)  
}  
  
outlier_detect<-get_outliers(df, "person_age")  
head(outlier_detect)
```

**Output:**

```
> outlier_detect<-get_outliers(df, "person_age")  
> head(outlier_detect)  
[1] 230 -22 -25 350 144 144
```

```
fix_outliers <- function(data, colname) {  
  if (!colname %in% names(data)) stop("Column not found.")  
  if (!is.numeric(data[[colname]])) stop("Column must be  
numeric.")  
  
  Q1 <- quantile(data[[colname]], 0.25, na.rm = TRUE)  
  Q3 <- quantile(data[[colname]], 0.75, na.rm = TRUE)  
  IQR <- Q3 - Q1  
  
  lower <- Q1 - 1.5 * IQR  
  upper <- Q3 + 1.5 * IQR
```

```

    is_outlier <- data[[colname]] < lower | data[[colname]] >
upper
    n_out <- sum(is_outlier, na.rm = TRUE)

    if (n_out > 0) {
      data[[colname]] <- pmin(pmax(data[[colname]], lower),
upper)
      message("Fixed ", n_out, " outliers in ", colname,
              " (capped to [", round(lower, 2), ", ",
round(upper, 2), "])")
    } else {
      message("No outliers to fix in ", colname)
    }

    return(data)
}
df_clean <- fix_outliers(df, "person_age")
head(as.data.frame(df_clean))

```

**Output:**

```

person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71945      0      RENT      35000      PERSONAL      16.02      0.49      3      561      0      Yes
2      21      female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      1      No
3      25      female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      0      Yes
4      23      female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      0      Yes
5      24      male      Master      66135      1      RENTT      35000      MEDICAL      14.27      0.53      4      586      0      Yes
6      23      female      High School      12951      0      OWN      2500      VENTURE      7.14      0.19      2      532      0      Yes
>

```

**Task 5: Normalization method for any continuous attribute:**

**Description:** In this task we apply normalization to a continuous column so that all its values are scaled into the range 0 to 1. We use the min–max normalization formula:

$$X_{norm} = \frac{x - \min(x)}{\max(x) - \min(x)}$$



Here we normalize the column `loan_amnt`. The smallest loan amount becomes 0, the largest loan amount becomes 1, and all other values are scaled proportionally between 0 and 1.

This step ensures that attributes measured on different scales can be compared fairly

```
df_nor <- df
normalize <- function(x) {
  return( (x - min(x)) / (max(x) - min(x)) )
}
df_nor$loan_amnt <- normalize(df_nor$loan_amnt)
head(df_nor)
```

### Output:

```
> head(df_nor)
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1         21         female          Master      71945         0          RENT 1.000000000  PERSONAL      16.02           0.49             3          561             0             1
2         21         female    High School      12282         0          OWN 0.000000000  EDUCATION      11.14           0.49             2          504             1             0
3         25         female    High School      12438         3    MORTGAGE 0.13235794  MEDICAL      12.87           0.49             3          635             0             1
4         23         female    Bachelor      79753         0          RENT 1.000000000  MEDICAL      15.23           0.44             2          675             0             1
5         24         male      Master      66135         1          RENT 1.000000000  MEDICAL      14.27           0.53             4          586             0             1
6         23         female    High School      12951         0          OWN 0.04411765  VENTURE       7.14           0.19             2          532             0             1
> |
```

## Task 6: Find and remove duplicate rows to reduce congestion:

### Description:

In this task we look for duplicate rows in the dataset and remove them to make the data cleaner and less congested. We use the `distinct()` function from the `dplyr` package. This function checks for duplicates based on a chosen column (here `person_education`) and keeps only the first occurrence, while the option `.keep_all = TRUE` ensures that all other columns in that row are also preserved.

```
df_unique <- df
df_unique <- distinct(df, person_education ,.keep_all = TRUE)
head(as.data.frame(df_unique))
```

**Output:**

```
> head(as.data.frame(df_unique))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1         21         Female         Master         71948             0             RENT         35000      PERSONAL         16.02             0.49             1          561             0             1
2         21         Female    High School         12282             0             OWN          1000      EDUCATION         11.14             0.49             2          504             1             0
3         23         Female    Bachelor         79753             0             RENT         35000      MEDICAL         15.23             0.44             2          675             0             1
4         21         Female    Associate         13113             0             OWN          4500    HOMEIMPROVEMENT         8.63             0.34             2          651             0             1
5         26         Female    Doctorate         56325             2             RENT         25000    DEBTCONSOLIDATION         11.86             0.44             4          690             0             1
>
```

**Description:****Task 7: Filter Data to find specific numerical value , row, categorical value:****Description:**

There are many ways to filter data in R, but here we are using the filter() function from the dplyr package. This allows us to select rows that match certain conditions. For example:

- Filtering numerical values such as loan\_amnt > 10000 to get only loans greater than 10,000.
- Filtering based on both missingness and value, such as keeping rows where person\_income is not missing and greater than 20,000.
- Filtering categorical values, such as selecting rows where person\_education is "Bachelor" or "Master".

```
filtered_data1 <- filter(df_clean, loan_amnt > 10000)
head(as.data.frame(filtered_data1))
filtered_data2 <- filter(df_clean, !is.na(person_income) &
person_income > 20000)
head(as.data.frame(filtered_data2))
filtered_data3 <- filter(df_clean, person_education %in%
c("Bachelor", "Master"))
head(as.data.frame(filtered_data3))
```

**Output:**

```
> head(as.data.frame(filtered_data1))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948.0      0      RENT      35000      PERSONAL      16.02      0.49      3      561      0      Yes
2      23      female      Bachelor      79753.0      0      RENT      35000      MEDICAL      15.23      0.44      2      675      0      Yes
3      24      male      Master      66135.0      1      RENT      35000      MEDICAL      14.27      0.53      4      586      0      Yes
4      22      female      Bachelor      149874.8      1      RENT      35000      EDUCATION      12.42      0.37      3      701      0      Yes
5      24      male      High School      95550.0      5      RENT      35000      MEDICAL      11.11      0.37      4      585      0      Yes
6      22      female      Bachelor      100684.0      3      RENT      35000      PERSONAL      8.90      0.35      2      544      0      Yes

> head(as.data.frame(filtered_data2))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948.0      0      RENT      35000      PERSONAL      16.02      0.49      3      561      0      Yes
2      23      female      Bachelor      79753.0      0      RENT      35000      MEDICAL      15.23      0.44      2      675      0      Yes
3      24      male      Master      66135.0      1      RENT      35000      MEDICAL      14.27      0.53      4      586      0      Yes
4      22      female      Bachelor      149874.8      1      RENT      35000      EDUCATION      12.42      0.37      3      701      0      Yes
5      24      male      High School      95550.0      5      RENT      35000      MEDICAL      11.11      0.37      4      585      0      Yes
6      22      female      Bachelor      100684.0      3      RENT      35000      PERSONAL      8.90      0.35      2      544      0      Yes

> head(as.data.frame(filtered_data3))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amnt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948.0      0      RENT      35000      PERSONAL      16.02      0.49      3      561      0      Yes
2      23      female      Bachelor      79753.0      0      RENT      35000      MEDICAL      15.23      0.44      2      675      0      Yes
3      24      male      Master      66135.0      1      RENT      35000      MEDICAL      14.27      0.53      4      586      0      Yes
4      22      female      Bachelor      149874.8      1      RENT      35000      EDUCATION      12.42      0.37      3      701      0      Yes
5      24      male      High School      95550.0      5      RENT      35000      MEDICAL      11.11      0.37      4      585      0      Yes
6      23      male      Bachelor      114860.0      3      RENT      35000      VENTURE      7.90      0.30      2      573      0      Yes
>
```

## Task8: Detect invalid value then remove this row or replace mean value:

### Description:

There are many ways to handle invalid/missing values; here we use two simple approaches:

- ❖ Remove rows with any missing values using `na.omit()` to get a clean subset and check how many rows remain.
- ❖ Impute (fill) missing values in the full dataset:
  - For numeric columns, replace NA with the mean of that column.
  - For categorical columns, replace NA with the mode (most frequent category).

After imputation, we verify the result by printing the count of remaining NAs per column group and previewing the cleaned data.

```
loans5<-Loan_Datas
loans5 <- na.omit(loans5)
print(nrow(loans5))
print(loans5)

loan5 <- Loan_Datas
loan3_clean <- na.omit(loan5)
print(nrow(loan3_clean))
loan3_clean
num_cols <- sapply(loan5, is.numeric)
```

```

cat_cols <- sapply(loan5, is.character)
loan5[num_cols] <- lapply(loan5[num_cols], function(x) {
  x[is.na(x)] <- mean(x, na.rm = TRUE)
  return(x)
})
loan5[cat_cols] <- lapply(loan5[cat_cols], function(x) {
  mode_val <- names(sort(table(x), decreasing = TRUE))[1]
  x[is.na(x)] <- mode_val
  return(x)
})
print(colSums(is.na(loan5[num_cols])))
print(colSums(is.na(loan5[cat_cols])))
print(loan5)

```

### Output:

```

> head(as.data.frame(loan5))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      Female      Master      71948      0      RENT      35000      PERSONAL      16.02      0.49      3      561      No      1
2      21      Female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      Yes      0
3      25      Female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      No      1
4      23      Female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      No      1
5      24      male      Master      66135      1      RENTT      35000      MEDICAL      14.27      0.53      4      586      No      1
6      21      Female      High School      12739      0      OWN      1600      VENTURE      14.74      0.13      3      640      No      1
> head(as.data.frame(loan5_clean))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      Female      Master      71948      0      RENT      35000      PERSONAL      16.02      0.49      3      561      No      1
2      21      Female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      Yes      0
3      25      Female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      No      1
4      23      Female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      No      1
5      24      male      Master      66135      1      RENTT      35000      MEDICAL      14.27      0.53      4      586      No      1
6      21      Female      High School      12739      0      OWN      1600      VENTURE      14.74      0.13      3      640      No      1
> head(as.data.frame(loan5))
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1 21.000000      Female      Master      71948      0      RENT      35000      PERSONAL      16.02      0.49      3      561      No      1
2 21.000000      Female      High School      12282      0      OWN      1000      EDUCATION      11.14      0.49      2      504      Yes      0
3 25.000000      Female      High School      12438      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      No      1
4 23.000000      Female      Bachelor      79753      0      RENT      35000      MEDICAL      15.23      0.44      2      675      No      1
5 24.000000      male      Master      66135      1      RENTT      35000      MEDICAL      14.27      0.53      4      586      No      1
6 26.913711      Female      High School      12951      0      OWN      2500      VENTURE      7.14      0.19      2      532      No      1
>

```

### Task 9. We can convert the imbalanced data set into the balanced data set

#### Description:

In our dataset, the target variable was highly imbalanced, so we applied three resampling techniques to balance it: undersampling, oversampling, and SMOTE.

- **Undersampling:** Here, we reduced the majority class to match the minority class. In the code, `N_under` was calculated as twice the size of the minority class, and the `ovun.sample()` function with `method = "under"` was used. This produced a balanced dataset but at the cost of discarding many majority samples, which reduced the overall dataset size.
- **Oversampling:** In this method, we increased the minority class to match the majority class. We set `N_over` as twice the size of the majority class and used `ovun.sample()` with

method = "over". This balanced the dataset by duplicating minority class samples, keeping all majority data but introducing the risk of overfitting due to repeated rows.

- **SMOTE (Synthetic Minority Oversampling Technique):** Unlike oversampling, SMOTE generates synthetic minority samples instead of duplicating them. In the code, categorical variables were converted into factors, incomplete rows were removed, and the ROSE() function was applied. This enriched the dataset with synthetic examples, helping the model generalize better while maintaining class balance.

Overall, undersampling removes data, oversampling duplicates data, and SMOTE creates synthetic data. These methods make the dataset more balanced, ensuring fairer and more accurate model training.

- **Undersampling:**

```
library(ROSE)
df$loan_status <- factor(df$loan_status, levels = c(0,1))
table(df$loan_status)
N_under <- 2 * min(table(df$loan_status))
set.seed(199)
under_df <- ovun.sample(loan_status ~ ., data = df, method = "under",
N = N_under, seed = 199)$data
table(under_df$loan_status)
head(under_df)
```

**Output:**

```
> table(df$loan_status)
0 1
76 125
> table(under_df$loan_status)
1 0
76 76
> head(under_df)
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1         21         female         Bachelor         35952             0             RENT         42000      PERSONAL         11.48             0.28             4             663             0             1
2         23         male         Bachelor         66599             0             RENT         25000      EDUCATION         11.36             0.38             3             653             0             1
3         23         male      High School         17755             1             RENT         4000      DEBTCONSOLIDATION 11.05             0.23             3             580             0             1
4         23         male         Bachelor         111553             3             RENT         30000      PERSONAL         15.23             0.27             3             648             0             1
5         26         female         Master         65880             2             RENT         25000      DEBTCONSOLIDATION 14.83             0.38             2             592             0             1
6         21         female         Bachelor         51922             3             RENT         25000      EDUCATION         12.73             0.48             3             673             0             1
>
```

- **Oversampling:**

```
library(ROSE)

df$loan_status <- factor(df$loan_status, levels = c(0,1))
table(df$loan_status)
N_over <- 2 * max(table(df$loan_status))
set.seed(199)
```

```
over_df <- ovun.sample(loan_status ~ ., data = df,
                      method = "over", N = N_over, seed = 199)$data
table(over_df$loan_status)
head(over_df)
```

**Output:**

```
> table(df$loan_status)
 0  1 
76 125
> table(over_df$loan_status)
 0  1 
125 125
> head(over_df)
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21      female      Master      71948.0      0      RENT      35000      PERSONAL      16.02      0.49      3      561      0      1
2      25      female      High School      12438.0      3      MORTGAGE      5500      MEDICAL      12.87      0.49      3      635      0      1
3      23      female      Bachelor      29753.0      0      RENT      35000      MEDICAL      15.23      0.44      2      675      0      1
4      24      male      Master      66135.0      1      RENT      35000      MEDICAL      14.27      0.53      4      586      0      1
5      23      female      High School      12951.0      0      OWN      2500      VENTURE      7.14      0.19      2      532      0      1
6      22      female      Bachelor      149874.8      1      RENT      35000      EDUCATION      12.42      0.37      3      701      0      1
```

- **Smote**

```
library(ROSE)
set.seed(199)
df_smote$previous_loan_defaults_on_file <-
factor(df_smote$previous_loan_defaults_on_file, levels = c(0,1))
df_smote[sapply(df_smote, is.character)] <-
lapply(df_smote[sapply(df_smote, is.character)], factor)
df_smote <- df_smote[complete.cases(df_smote), ]
table(df_smote$previous_loan_defaults_on_file)
rose_df <- ROSE(previous_loan_defaults_on_file ~ ., data = df_smote,
N = 2000, p = 0.5)$data
table(rose_df$previous_loan_defaults_on_file)
head(rose_df)
```

**Output:**

```
> table(df_smote$previous_loan_defaults_on_file)
 0  1 
150 51
> table(rose_df$previous_loan_defaults_on_file)
 0  1 
1004 996
> head(rose_df)
  person_age person_gender person_education person_income person_emp_exp person_home_ownership loan_amt loan_intent loan_int_rate loan_percent_income cb_person_cred_hist_length credit_score previous_loan_defaults_on_file loan_status
1      21.22442      female      Bachelor      26171.050      4.1156402      RENT      32192.34      EDUCATION      12.75038      0.3815402      1.418469      708.3967      0      Yes
2      29.10654      female      High School      297281.332      10.1892669      RENT      24690.35      HOMEIMPROVEMENT      13.98243      0.4329616      1.981410      598.3117      0      Yes
3      21.66972      female      Master      334125.360      7.1406131      RENT      29846.81      PERSONAL      16.06708      0.4950736      1.925310      547.3351      0      Yes
4      25.49819      female      Bachelor      -70203.018      -13.4609624      RENT      20956.51      DEBTCONSOLIDATION      12.79445      0.5623364      3.718365      650.1541      0      No
5      25.10655      male      Bachelor      92247.763      2.4591532      RENT      15756.08      EDUCATION      13.88743      0.2648101      2.709815      573.3217      0      Yes
6      26.28129      male      Bachelor      -4518.721      -0.8790874      RENT      22708.11      EDUCATION      16.05170      0.2265051      2.175995      625.9735      0      Yes
```

**Task 10: Split the dataset for Training and Testing , 70% row for Training data and 30% for Testing data:**

**Description:**

In this task we divide the dataset into two parts: 70% for training and 30% for testing. The function `initial_split()` from the `rsample` package is used to create the split. From this split object, the `training()` function extracts the training set, and the `testing()` function extracts the testing set.

The training set is used to fit and build the model, while the testing set is kept aside to check how well the model performs on new, unseen data. Finally, the `dim()` function shows the number of rows and columns in each set to confirm that the split was done correctly (about 70% of rows in training and 30% in testing).

```
set.seed(123)
split <- initial_split(df, prop = 0.7)
train_data <- training(split)
test_data <- testing(split)
dim(train_data)
dim(test_data)
```

Output:

```
> dim(train_data)
[1] 140 14
> dim(test_data)
[1] 61 14
> |
```

**Description:****11. statistics and interpret the results for the following numerical:**

variables between two target classes (`loan_status = 1` and `loan_status = 0`)

➤ `person_age`

➤ Person Income:

**Description:**

In this task we calculate summary statistics of two numerical variables — `person_age` and `person_income` — separately for the two target classes of `loan_status` (0 and 1).

We use the `aggregate()` function with a custom summary that computes the mean, median, standard deviation, minimum, and maximum for each variable within each class. This produces grouped summaries:

- For `loan_status = 0` (non-default customers)
- For `loan_status = 1` (default customers)

The results help us interpret how age and income differ between the two classes. For example, we can see whether defaulters tend to be younger or older, and whether their average income is higher or lower compared to non-defaulters. This comparison gives useful insights into which demographic or financial factors may be linked to loan outcomes.

```
age_summary <- aggregate(person_age ~ loan_status, data = df,
                          FUN = function(x) c(mean = mean(x), median =
median(x),
                                              sd = sd(x), min =
min(x), max = max(x)))
age_summary <- do.call(data.frame, age_summary)

# Create summary table for Income
income_summary <- aggregate(person_income ~ loan_status, data = df,
                             FUN = function(x) c(mean = mean(x),
median = median(x),
                                              sd = sd(x), min =
min(x), max = max(x)))
income_summary <- do.call(data.frame, income_summary)

print(age_summary)
```

**Output:**

```
  loan_status person_age.mean person_age.median person_age.sd person_age.min person_age.max
1         No    29.10526         24      30.88757         -22         230
2         Yes    25.45600         23      29.61895         -25         350
> |
```



```
print(income_summary)
```

**Output:**

```
 loan_status person_income.mean person_income.median person_income.sd person_income.min person_income.max
1          No      232460.21      249174          95317.19         12282         368115
2          Yes       99662.79       72608        279363.03         12438        3138998
>
```

### Task 12 .Compare average credit\_score between customers with loan\_status 1 and with loan\_status 0:

#### Description:

To compare the average credit score between customers with different loan statuses, we applied the `aggregate()` function in R. The dataset was grouped by `loan_status`, where 0 represents customers without default (good status) and 1 represents customers with default (bad status). The mean credit score was then calculated for each group, rounded to two decimal places, and missing values were ignored (`na.rm=TRUE`).

The results show that customers with `loan_status = 0` have an average credit score of 630.09, while customers with `loan_status = 1` have an average credit score of 627.55. This indicates that, on average, customers without default (0) have slightly higher credit scores than those with default (1).

```
aggregate(credit_score ~ loan_status, data = df, FUN = function(x)
round(mean(x, na.rm=TRUE), 2))
```

**Output:**

```
> aggregate(credit_score ~ loan_status, data = df, FUN = function(x) round(mean(x, na.rm=TRUE), 2))
  loan_status credit_score
1          0         630.09
2          1         627.55
>
```

**Description:**

Here group by function used for find out all value such as mean, median, mode, max, min for each group. Here has 0 and 1 group. Summarise function use computing statistics like mean, median, sum, count etc.

**13. Compare spread in person\_emp\_exp for customers with different levels of person\_education:****Description:**

In this task we are analyzing how the variable person\_emp\_exp (employment experience in years) is distributed across different categories of person\_education (such as High School, Bachelor, Master, etc.).

We created a custom function compare\_spread(), which groups the dataset by the specified categorical variable (person\_education) and then calculates descriptive spread measures for the numerical variable (person\_emp\_exp). Specifically, it computes:

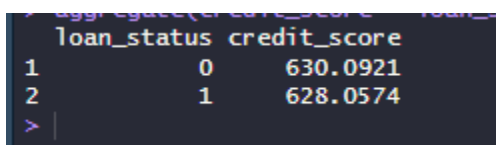
- Count: Number of records in each education level.
- Mean: The average employment experience.
- Standard Deviation (SD): How much the experience varies within that group.
- Minimum and Maximum: The range of values observed.
- Interquartile Range (IQR): The spread of the middle 50% of the data.

By comparing these statistics, we can see whether people with higher education levels tend to have more or less work experience, and whether the variation (spread) in experience differs

between groups. For example, graduates may have higher average experience but also larger variation compared to high school educated individuals.

**CODE:**

```
compare_spread <- function(data, group_col, value_col) {  
  library(dplyr)  
  
  if (!group_col %in% names(data)) stop("Group column not found.")  
  if (!value_col %in% names(data)) stop("Value column not found.")  
  
  data %>%  
    group_by(.data[[group_col]]) %>%  
    summarise(  
      count = n(),  
      mean = mean(.data[[value_col]], na.rm = TRUE),  
      sd = sd(.data[[value_col]], na.rm = TRUE),  
      min = min(.data[[value_col]], na.rm = TRUE),  
      max = max(.data[[value_col]], na.rm = TRUE),  
      IQR = IQR(.data[[value_col]], na.rm = TRUE),  
      .groups = "drop"  
    )  
}  
aggregate(credit_score ~ loan_status, data = data, mean, na.rm =  
TRUE)
```

**Output:**

```
> aggregate(credit_score ~ loan_status, data = data, mean, na.rm = TRUE)  
  loan_status credit_score  
1          0    630.0921  
2          1    628.0574  
> |
```

```
compare <- compare_spread(df_clean, "person_education",  
"person_emp_exp")  
head(as.data.frame(compare))
```

**Output:**

```
> head(as.data.frame(compare))
  person_education count      mean      sd min max IQR
1      Associate    46 3.760870 17.724039   0 121  2.0
2      Bachelor    73 3.534247 14.533565   0 125  3.0
3      Doctorate     1 2.000000      NA    2   2  0.0
4   High School    58 1.413793  1.797020   0   7  3.0
5         Master    23 1.739130  1.888178   0   6  2.5
> |
```

**Project Code**

```
library(readxl)
library(modeest)
library(naniar)
library(dplyr)
library(rsample)

data <- read_excel("D:/data science project-
mid/data/Midterm_Dataset_Section(A).xlsx")

missing_rows <- data[!complete.cases(data), ]
head(missing_rows)

drop_data <- na.omit(data);
head(drop_data);

df_mean <- data
df_mean$person_income[is.na(df_mean$person_income)] <-
mean(df_mean$person_income, na.rm = TRUE);
head(df_mean)

df_median <- df_mean
df_median$person_age[is.na(df_median$person_age)] <-
median(df_median$person_age, na.rm = TRUE);
head(df_median)

df_mode <- df_median
mode_val <- mlv(df_mode$loan_status, method = "mfv", na.rm
= TRUE);
df_mode$loan_status[is.na(df_mode$loan_status)] <- mode_val
head(as.data.frame(df_mode))

mode_val <- mlv(df_mode$person_gender, method = "mfv",
na.rm = TRUE);
df_mode$person_gender[is.na(df_mode$person_gender)] <-
mode_val
```

```
mode_val <- mlv(df_mode$person_education, method = "mfv",
na.rm = TRUE);
df_mode$person_education[is.na(df_mode$person_education)]
<- mode_val

print(sum(is.na(df_mode)))

vis_miss(data)

df_mode$loan_status <- ifelse(df_mode$loan_status == 1,
"Yes", "No");
head(as.data.frame(df_mode))

df_mode$previous_loan_defaults_on_file <-
ifelse(df_mode$previous_loan_defaults_on_file == "Yes", 1,
0);
head(as.data.frame(df_mode))
df <- df_mode

get_outliers <- function(data, colname) {
  if (!colname %in% names(data)) stop("Column not found.")
  if (!is.numeric(data[[colname]])) stop("Column must be
numeric.")

  Q1 <- quantile(data[[colname]], 0.25, na.rm = TRUE)
  Q3 <- quantile(data[[colname]], 0.75, na.rm = TRUE)
  IQR <- Q3 - Q1

  lower <- Q1 - 1.5 * IQR
  upper <- Q3 + 1.5 * IQR
```

```
    outliers <- data[[colname]][data[[colname]] < lower |
data[[colname]] > upper]
    return(outliers)
}

fix_outliers <- function(data, colname) {
  if (!colname %in% names(data)) stop("Column not found.")
  if (!is.numeric(data[[colname]])) stop("Column must be
numeric.")

  Q1 <- quantile(data[[colname]], 0.25, na.rm = TRUE)
  Q3 <- quantile(data[[colname]], 0.75, na.rm = TRUE)
  IQR <- Q3 - Q1

  lower <- Q1 - 1.5 * IQR
  upper <- Q3 + 1.5 * IQR

  is_outlier <- data[[colname]] < lower | data[[colname]] >
upper
  n_out <- sum(is_outlier, na.rm = TRUE)

  if (n_out > 0) {
    data[[colname]] <- pmin(pmax(data[[colname]], lower),
upper)
  }

  return(data)
}

df <- df_mode
outlier_detect <- get_outliers(df, "person_age")
head(outlier_detect)

df_clean <- fix_outliers(df, "person_age")
head(as.data.frame(df_clean))

df_nor <- df
```

```
normalize <- function(x) {  
  return((x - min(x)) / (max(x) - min(x)))  
}  
df_nor$loan_amnt <- normalize(df_nor$loan_amnt)  
head(df_nor)  
  
df_unique <- df  
df_unique <- distinct(df, person_education, .keep_all =  
TRUE)  
head(as.data.frame(df_unique))  
  
filtered_data1 <- filter(df_clean, loan_amnt > 10000)  
head(as.data.frame(filtered_data1))  
  
filtered_data2 <- filter(df_clean, !is.na(person_income) &  
person_income > 20000)  
head(as.data.frame(filtered_data2))  
  
filtered_data3 <- filter(df_clean, person_education %in%  
c("Bachelor", "Master"))  
head(as.data.frame(filtered_data3))  
  
loans5 <- data  
loans5 <- na.omit(loans5)  
print(nrow(loans5))  
head(as.data.frame(loans5))  
  
loan5 <- data  
loan3_clean <- na.omit(loan5)  
print(nrow(loan3_clean))  
head(as.data.frame(loan3_clean))  
  
num_cols <- sapply(loan5, is.numeric)
```



```
cat_cols <- sapply(loan5, is.character)
loan5[num_cols] <- lapply(loan5[num_cols], function(x) {
  x[is.na(x)] <- mean(x, na.rm = TRUE)
  return(x)
})
loan5[cat_cols] <- lapply(loan5[cat_cols], function(x) {
  mode_val <- names(sort(table(x), decreasing = TRUE))[1]
  x[is.na(x)] <- mode_val
  return(x)
})
print(colSums(is.na(loan5[num_cols])))
print(colSums(is.na(loan5[cat_cols])))
head(as.data.frame(loan5))

df$loan_status <- ifelse(df$loan_status == "Yes", 1, 0);
under_df <- df

library(ROSE)
df$loan_status <- factor(df$loan_status, levels = c(0,1))
table(df$loan_status)
N_under <- 2 * min(table(df$loan_status))
set.seed(199)
under_df <- ovun.sample(loan_status ~ ., data = df, method
= "under", N = N_under, seed = 199)$data
table(under_df$loan_status)
head(under_df)

df$loan_status <- factor(df$loan_status, levels = c(0,1))
table(df$loan_status)
N_over <- 2 * max(table(df$loan_status))
set.seed(199)
over_df <- ovun.sample(loan_status ~ ., data = df, method =
"over", N = N_over, seed = 199)$data
table(over_df$loan_status)
head(over_df)

df_smote <- df_clean
```

```
set.seed(199)
df_smote$previous_loan_defaults_on_file <-
factor(df_smote$previous_loan_defaults_on_file, levels =
c(0,1))
df_smote[sapply(df_smote, is.character)] <-
lapply(df_smote[sapply(df_smote, is.character)], factor)
df_smote <- df_smote[complete.cases(df_smote), ]

table(df_smote$previous_loan_defaults_on_file)
rose_df <- ROSE(previous_loan_defaults_on_file ~ ., data =
df_smote, N = 2000, p = 0.5)$data
table(rose_df$previous_loan_defaults_on_file)
head(rose_df)

orig <- names(df_clean)
cat_vars <- setdiff(orig[sapply(df_smote, function(x)
is.factor(x) || is.character(x))],
                    "previous_loan_defaults_on_file")

pretty_df <- df_clean
for (v in cat_vars) {
  d <- grep(paste0("^", v, "_"), names(pretty_df), value =
TRUE)
  if (length(d)) {
    M <- as.matrix(pretty_df[, d, drop = FALSE])
    lv <- sub(paste0("^", v, "_"), "", d)
    pretty_df[[v]] <- factor(lv[max.col(M, ties.method =
"first")],
                           levels =
levels(df_smote[[v]]))
    pretty_df[d] <- NULL
  }
}

pretty_df <- pretty_df[, orig, drop = FALSE]
head(as.data.frame(pretty_df))
```

```
set.seed(123)
split <- initial_split(df, prop = 0.7)
train_data <- training(split)
test_data <- testing(split)
dim(train_data)
dim(test_data)

age_summary <- aggregate(person_age ~ loan_status, data =
df,
                        FUN = function(x) c(mean =
mean(x), median = median(x),
                        sd = sd(x),
min = min(x), max = max(x)))
age_summary <- do.call(data.frame, age_summary)

income_summary <- aggregate(person_income ~ loan_status,
data = df,
                        FUN = function(x) c(mean =
mean(x), median = median(x),
                        sd = sd(x),
min = min(x), max = max(x)))
income_summary <- do.call(data.frame, income_summary)

print(age_summary)
print(income_summary)

aggregate(credit_score ~ loan_status, data = df, FUN =
function(x) round(mean(x, na.rm=TRUE), 2))

compare_spread <- function(data, group_col, value_col) {
  library(dplyr)
```

```
if (!group_col %in% names(data)) stop("Group column not
found.")
if (!value_col %in% names(data)) stop("Value column not
found.")

data %>%
  group_by(.data[[group_col]]) %>%
  summarise(
    count = n(),
    mean = mean(.data[[value_col]], na.rm = TRUE),
    sd = sd(.data[[value_col]], na.rm = TRUE),
    min = min(.data[[value_col]], na.rm = TRUE),
    max = max(.data[[value_col]], na.rm = TRUE),
    IQR = IQR(.data[[value_col]], na.rm = TRUE),
    .groups = "drop"
  )
}

aggregate(credit_score ~ loan_status, data = data , mean,
na.rm = TRUE)
compare <- compare_spread(df_clean, "person_education",
"person_emp_exp")
head(as.data.frame(compare))
```